

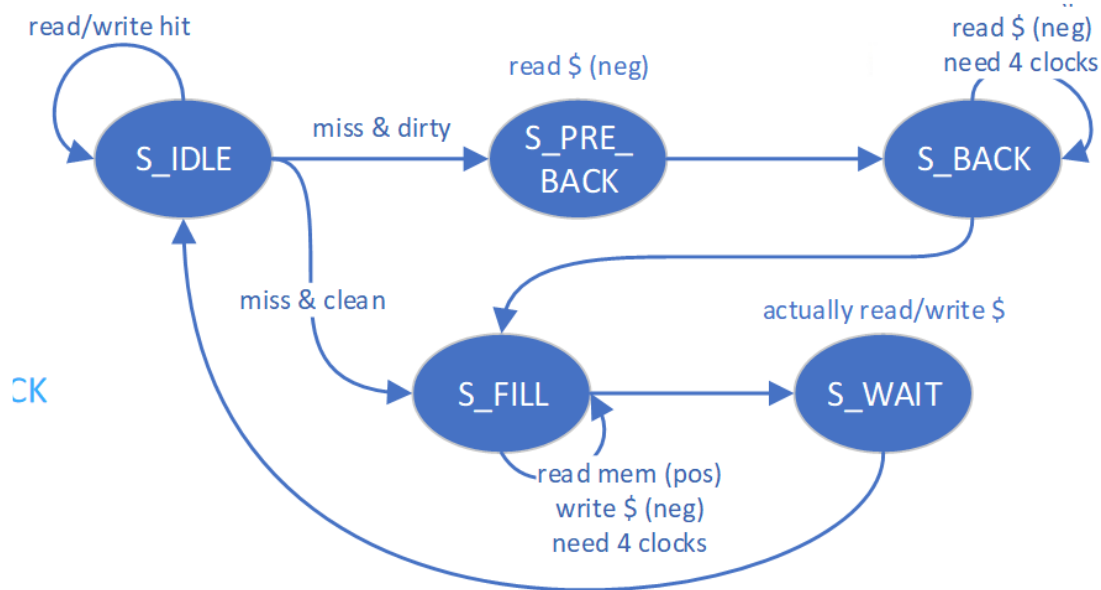
# Lab4: Pipelined CPU with Cache

3220106039 李瀚轩

## 1. Experimental Steps

### 1.1 Complete the CMU Module

According to the PPT, the key to the CMU module is completing the state machine setup, based on the diagram below:



After this, we can complete the framework code:

```
module cmu (  
    // CPU side  
    input clk,  
    input rst,  
    input [31:0] addr_rw,  
    input en_r,  
    input en_w,  
    input [2:0] u_b_h_w,  
    input [31:0] data_w,  
    output [31:0] data_r,  
    output stall,  
  
    // mem side  
    output reg mem_cs_o = 0,  
    output reg mem_we_o = 0,  
    output reg [31:0] mem_addr_o = 0,  
    input [31:0] mem_data_i,  
    output [31:0] mem_data_o,  
    input mem_ack_i,  
  
    // debug info
```

```

        output [2:0] cmu_state
    );

    `include "addr_define.vh"

    reg [ADDR_BITS-1:0] cache_addr = 0;
    reg cache_load = 0;
    reg cache_store = 0;
    reg cache_edit = 0;
    reg [2:0] cache_u_b_h_w = 0;
    reg [WORD_BITS-1:0] cache_din = 0;
    wire cache_hit;
    wire [WORD_BITS-1:0] cache_dout;
    wire cache_valid;
    wire cache_dirty;
    wire [TAG_BITS-1:0] cache_tag;

    cache CACHE (
        .clk(~clk),
        .rst(rst),
        .addr(cache_addr),
        .load(cache_load),
        .store(cache_store),
        .edit(cache_edit),
        .invalid(1'b0),
        .u_b_h_w(cache_u_b_h_w),
        .din(cache_din),
        .hit(cache_hit),
        .dout(cache_dout),
        .valid(cache_valid),
        .dirty(cache_dirty),
        .tag(cache_tag)
    );

    localparam
        S_IDLE = 0,
        S_PRE_BACK = 1,
        S_BACK = 2,
        S_FILL = 3,
        S_WAIT = 4;

    reg [2:0] state = 0;
    reg [2:0] next_state = 0;
    reg [ELEMENT_WORDS_WIDTH-1:0] word_count = 0;
    reg [ELEMENT_WORDS_WIDTH-1:0] next_word_count = 0;
    assign cmu_state = state;

    always @ (posedge clk) begin
        if (rst) begin
            state <= S_IDLE;
            word_count <= 2'b00;
        end
        else begin
            state <= next_state;

```

```

        word_count ≤ next_word_count;
    end
end

// state ctrl
always @ (*) begin
    if (rst) begin
        next_state = S_IDLE;
        next_word_count = 2'b00;
    end
    else begin
        case (state)
            S_IDLE: begin
                if (en_r || en_w) begin
                    if (cache_hit)
                        next_state = S_IDLE;
                    else if (cache_valid && cache_dirty)
                        next_state = S_PRE_BACK;
                    else
                        next_state = S_FILL;
                end
                next_word_count = 2'b00;
            end

            S_PRE_BACK: begin
                next_state = S_BACK;
                next_word_count = 2'b00;
            end

            S_BACK: begin
                if (mem_ack_i && word_count ==
{ELEMENT_WORDS_WIDTH{1'b1}}) // 2'b11 in default case
                    next_state = S_FILL;
                else
                    next_state = S_BACK;

                if (mem_ack_i)
                    next_word_count = word_count + 2'b01;
                else
                    next_word_count = word_count;
            end

            S_FILL: begin
                if (mem_ack_i && word_count ==
{ELEMENT_WORDS_WIDTH{1'b1}})
                    next_state = S_WAIT;
                else
                    next_state = S_FILL;

                if (mem_ack_i)
                    next_word_count = word_count + 2'b01;
                else
                    next_word_count = word_count;
            end
        endcase
    end
end

```

```

        S_WAIT: begin
            next_state = S_IDLE;
            next_word_count = 2'b00;
        end
    endcase
end
end

// cache ctrl
always @ (*) begin
    case(state)
        S_IDLE, S_WAIT: begin
            cache_addr = addr_rw;
            cache_load = en_r;
            cache_edit = en_w;
            cache_store = 1'b0;
            cache_u_b_h_w = u_b_h_w;
            cache_din = data_w;
        end
        S_BACK, S_PRE_BACK: begin
            cache_addr = {addr_rw[ADDR_BITS-1:BLOCK_WIDTH],
next_word_count, {ELEMENT_WORDS_WIDTH{1'b0}}};
            cache_load = 1'b0;
            cache_edit = 1'b0;
            cache_store = 1'b0;
            cache_u_b_h_w = 3'b010;
            cache_din = 32'b0;
        end
        S_FILL: begin
            cache_addr = {addr_rw[ADDR_BITS-1:BLOCK_WIDTH], word_count,
{ELEMENT_WORDS_WIDTH{1'b0}}};
            cache_load = 1'b0;
            cache_edit = 1'b0;
            cache_store = mem_ack_i;
            cache_u_b_h_w = 3'b010;
            cache_din = mem_data_i;
        end
    endcase
end
assign data_r = cache_dout;

// mem ctrl
always @ (*) begin
    case (next_state)
        S_IDLE, S_PRE_BACK, S_WAIT: begin
            mem_cs_o = 1'b0;
            mem_we_o = 1'b0;
            mem_addr_o = 32'b0;
        end
        S_BACK: begin
            mem_cs_o = 1'b1;
            mem_we_o = 1'b1;

```

```

        mem_addr_o = {cache_tag, addr_rw[ADDR_BITS-TAG_BITS-
1:BLOCK_WIDTH], next_word_count, {ELEMENT_WORDS_WIDTH{1'b0}}};
    end

    S_FILL: begin
        mem_cs_o = 1'b1;
        mem_we_o = 1'b0;
        mem_addr_o = {addr_rw[ADDR_BITS-1:BLOCK_WIDTH], next_word_count,
{ELEMENT_WORDS_WIDTH{1'b0}}};
    end
endcase
end
assign mem_data_o = cache_dout;

    assign stall = next_state ≠ S_IDLE;

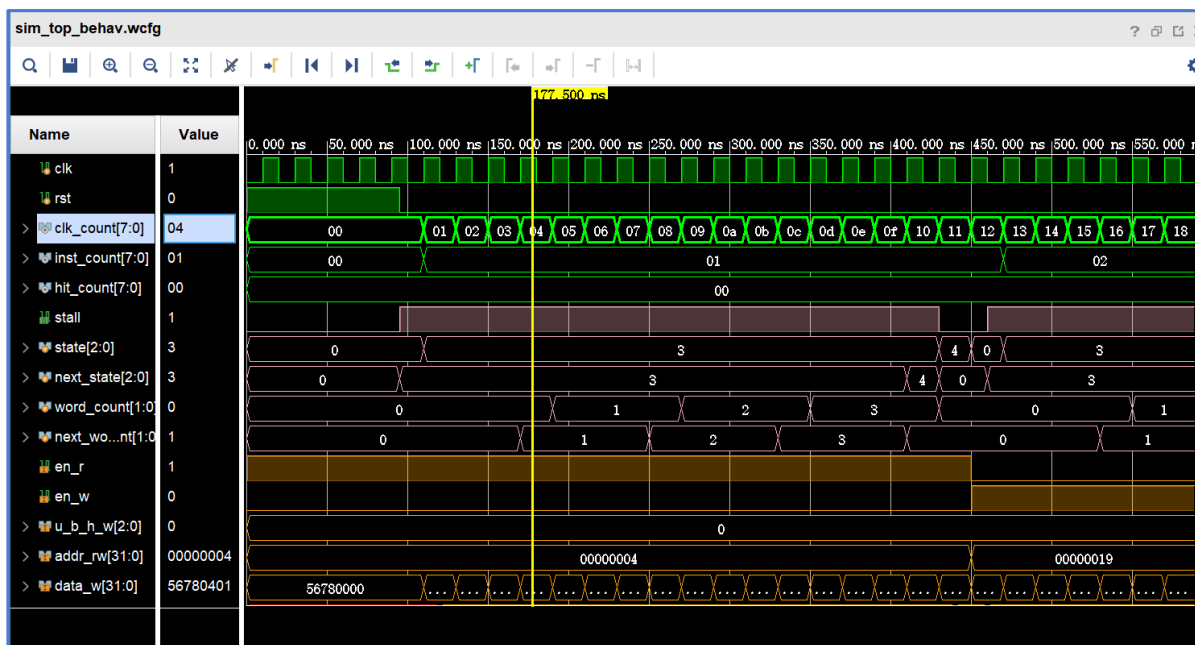
endmodule

```

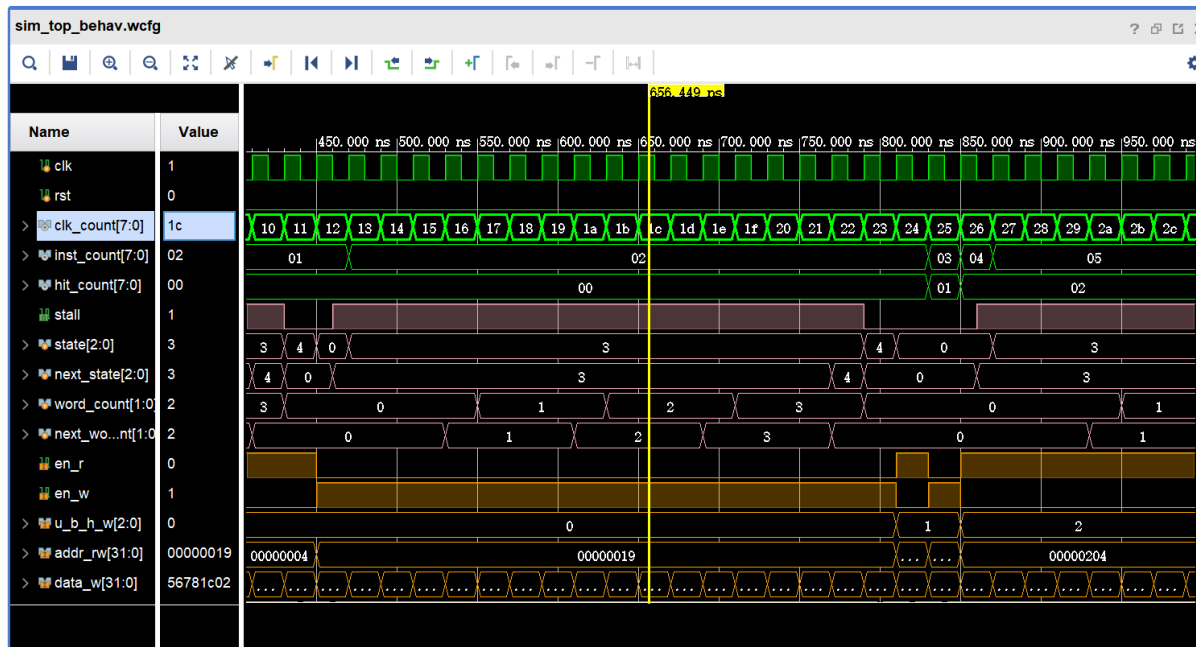
## 1.2 Completing the Cache Module

Here, we can simply bring over the code from lab3 without further explanation.

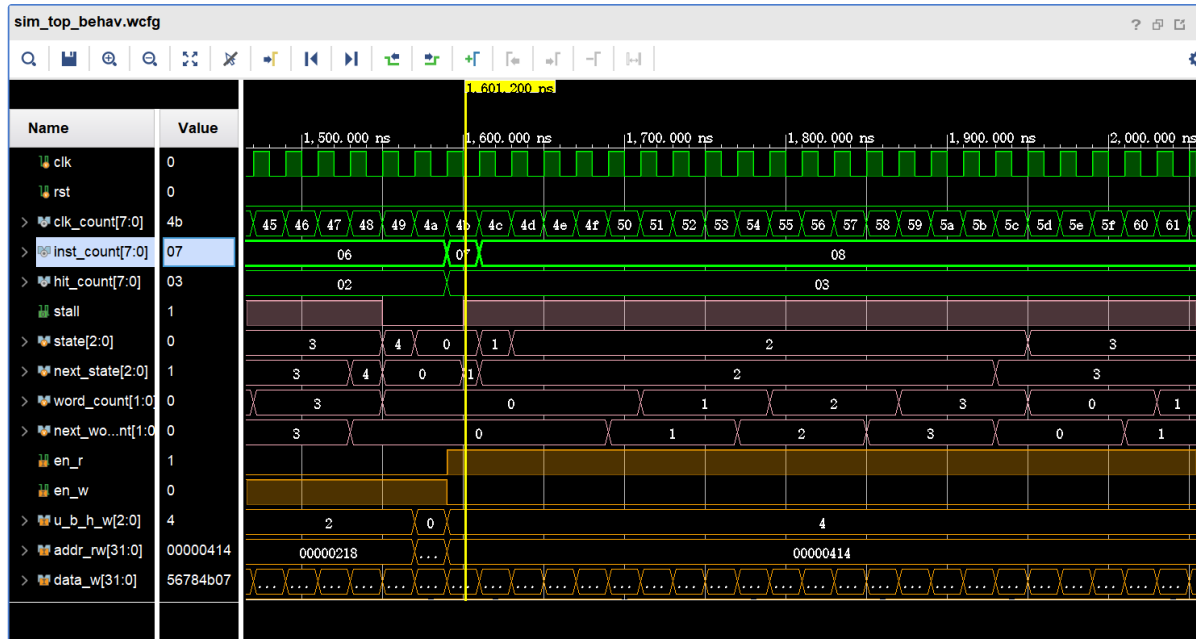
## 2. Evaluation



The instruction at 0x4 is a load instruction. Since the cache is empty, a read miss occurs. As a result, we can observe 0x8 and 0xC changing sequentially (with mem\_addr\_o going from 0x0 to 0x4 and cache\_store=1), which loads the corresponding data into the cache. Here, you can see that it takes 4\*4=16 cycles (with clk\_count changing from 0 to 0x10).



For the third and fourth instructions, a cache hit occurs. At 830ns and 850ns, you can see that one is a read hit and the other is a write hit, each taking only one cycle.



For the eighth instruction, a miss occurs, and a write-back is required. At 1610ns, the state changes from 0 to 1 (indicating `S_PRE_BACK`), and in the next cycle, it changes to 2 (indicating `S_BACK`). During this phase, the cache data is written back, and then it moves into the `S_FILL` phase, each requiring  $4 \times 4 = 16$  cycles.